




Explore and Understand
the Amount of Fabrication
of a Topic in Twitter

BotBoosted



Brian Balagot
brian.balagot@gmail.com
<https://www.linkedin.com/in/briancbalagot>
<https://github.com/brityboy>





NewsVideoBusinessNewsbreakMovePHViewsLife & StyleEntertainmentSportsTechBr

SECTIONS

 HOME

 SEARCH

The New York Times

INVESTIGATIVE

Fake accounts, manufactured reality on social media

Suspicious and fake Facebook accounts share common practices, forming a nest that spreads and repeats lies and propaganda



Chay F. Hofileña

@chayhoflena

Published 10:04 AM, October 09, 2016

Updated 9:55 AM, October 15, 2016

56

17K

Twitter



ILLUSTRATION BY EMIL MERCADO

Fake News in U.S. Election? Elsewhere, That's Nothing New

By PAUL MOZUR and MARK SCOTT NOV. 17, 2016



An Internet cafe in Manila. The Philippines is one of Facebook's fastest-growing markets.
Joe Dirietro/Agence France Presse Getty Images

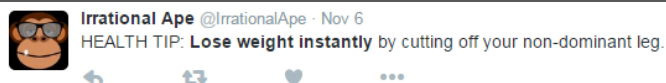
An **app** that can help me **understand how much** about a topic on an online social network **is fabricated or not,** and **what the conversations are** about

To explore and understand the amount of fabrication on Twitter...

*A tweet is **fabricated** if it was not made by a bona-fide human or genuine news source*


Search: lose weight instantly

REAL

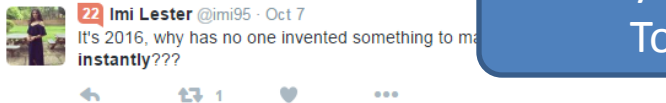


#1: Lightweight Supervised Classifier


REAL



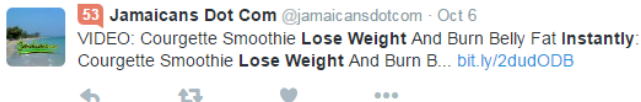
REAL



FAKE



REAL



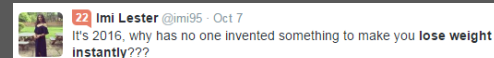
"think invent
try lose
weight"

100% real

#2: Dynamic Lightweight
Topic Modeler


"learn shed
pounds burn
belly fat
smoothie"

50% real



REAL

#3: Corpus Summarizer

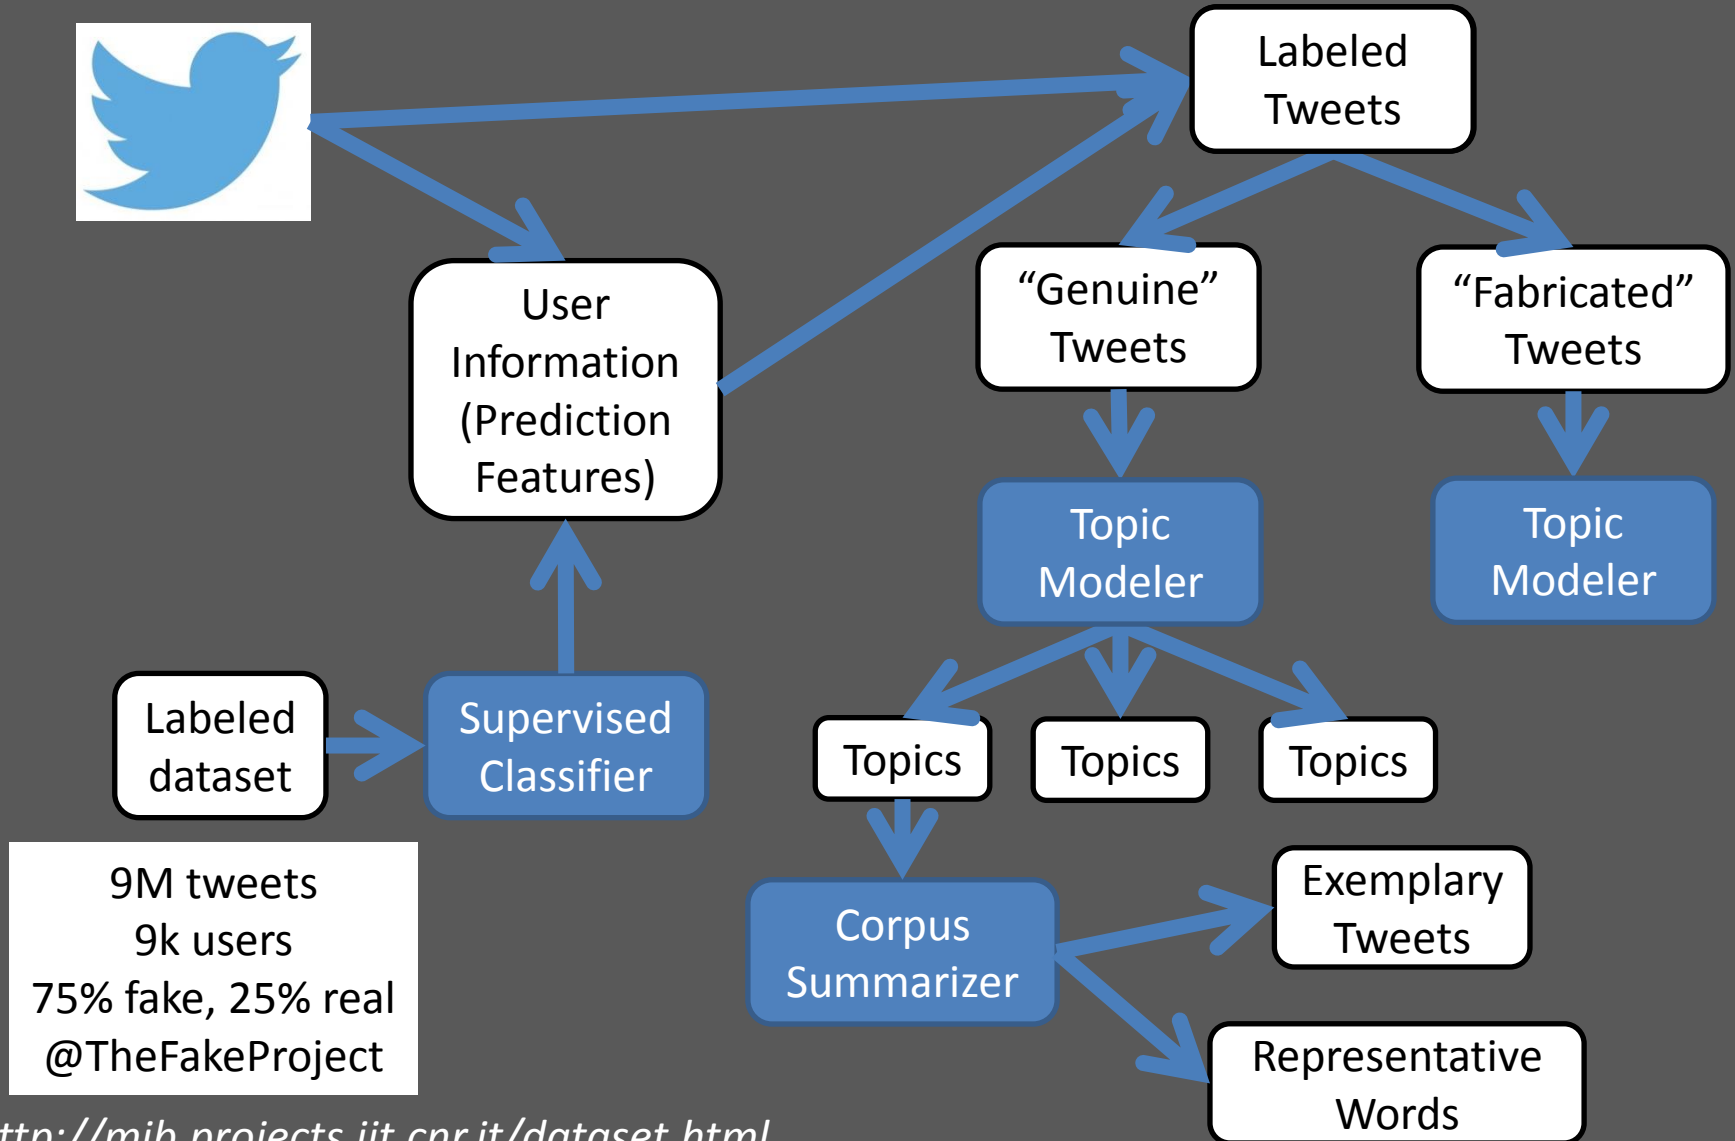


FAKE



REAL

My Data Strategy / Pipeline



1: Lightweight Supervised Classifier

The GOAL: Classify a user, based on their most recent tweet, as real or fake



Account History: followers_count,
friends_count, total_tweets,
total_likes

Relative_Volume:
likes_friends,
tweets_friends

Behavior Rate: likes_day,
tweets_day, friends_day,
likes_friends_day

10FCV
Acc: 97%

Account
History

Random
Forest

10FCV
Acc: 95%

Behavior Rate

Random
Forest

Account History

Behavior Rate

History Pred %

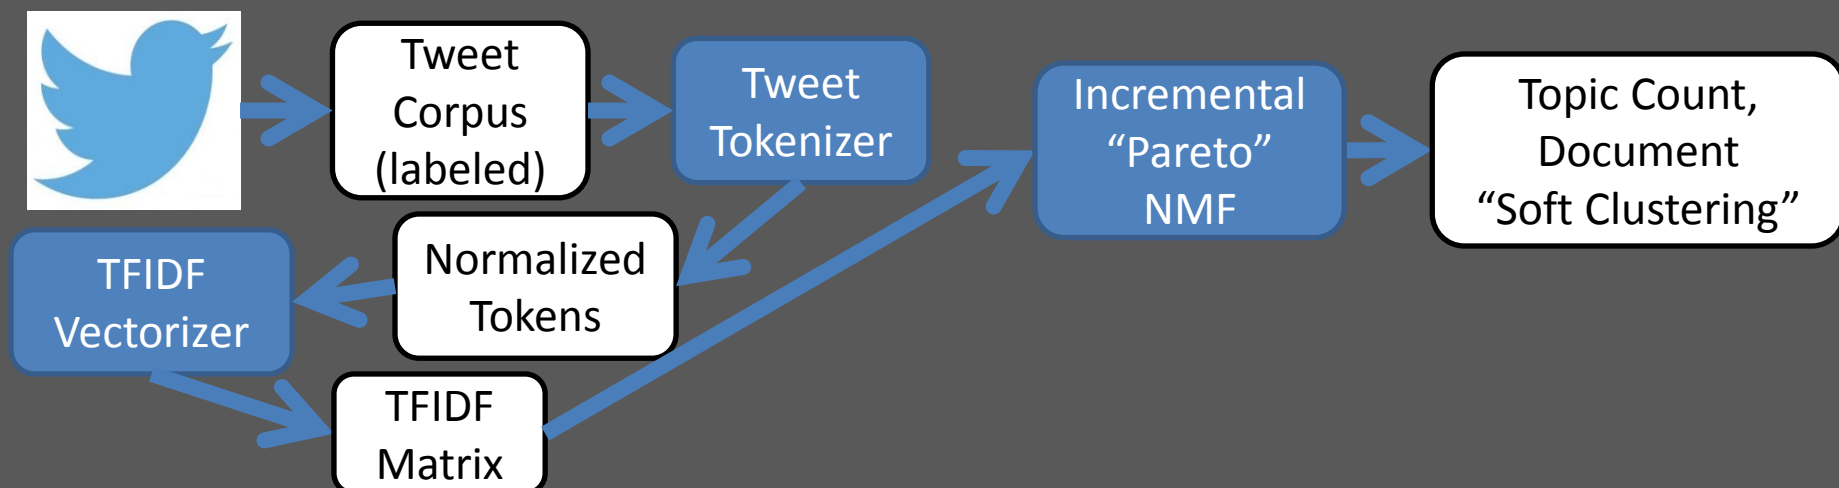
Rate Pred %

Random
Forest

10FCV
Acc: 98%

2: Dynamic Lightweight Topic Modeler

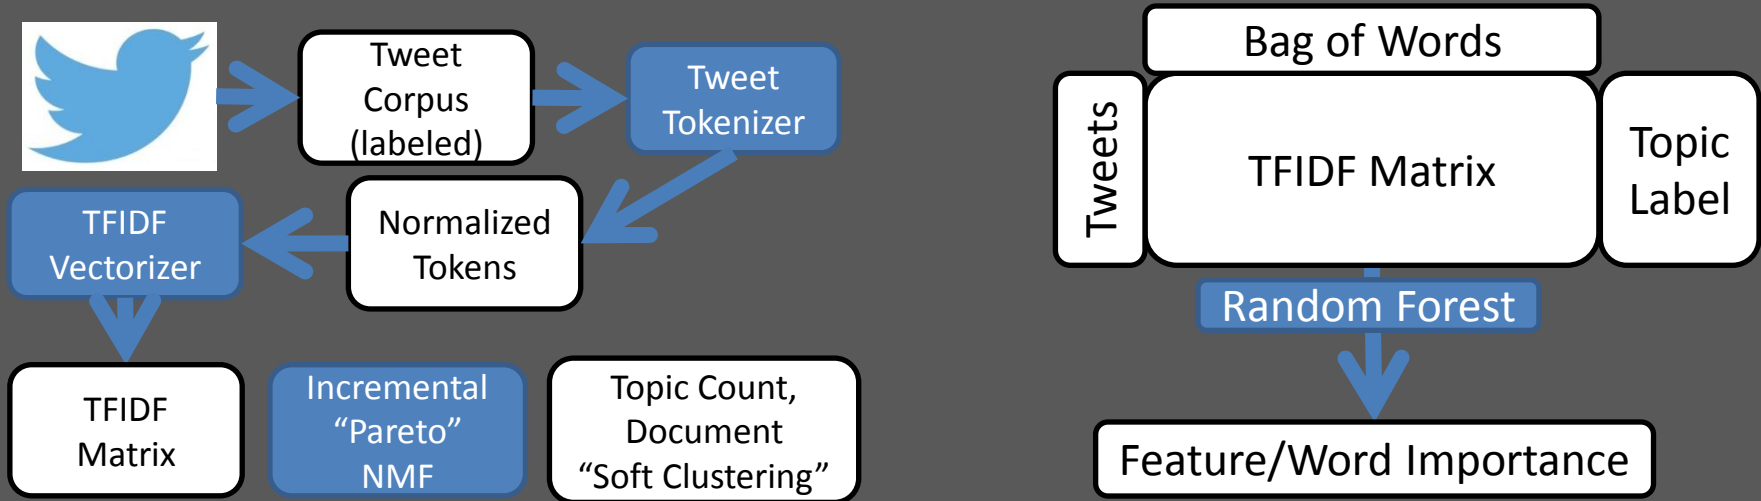
The GOAL: *Heuristically determine the topic count of a corpus, on the fly*



```
initializing evaluation...
extracting 2 topics...
[9361 2543] is the topic distribution
extracting 4 topics...
[6683 2719 1933 569] is the topic distribution
extracting 6 topics...
[4284 2266 2111 1616 1105 522] is the topic distribution
extracting 8 topics...
[3289 2244 1918 1592 1093 1061 517 190] is the topic distribution
extracting 10 topics...
[2232 1721 1634 1583 1162 1135 1116 515 468 338] is the topic distribution
heuristic topic count is 8
```

3: Corpus Summarizer

The GOAL: *extract the most important tweet that captures the subtopic*



A tweet is important if it is made up of important words

A word is important if it is helpful in "bucketing" tweets into topics

$$\text{Word Importance} = \text{Term Frequency} \times \text{Inverse Document Frequency} \times \text{Feature/Word Importance}$$


$$\text{Tweet Importance} = \text{Highest Total Word Importance}$$

Tweets by Topic, by Real/Fake, for: lose weight now

There are 423 fake tweets and 1394 real tweets

Tweets by Topic, by Real/Fake, for: lose weight now

There are 423 fake tweets and 1394 real tweets



want
right
weight
lose weight
lose

url_url_
weight url_
lose weight
weight

user_
weight
lose
lose weight
user_user_

hash_
hash_
url_
hash_url_
user_

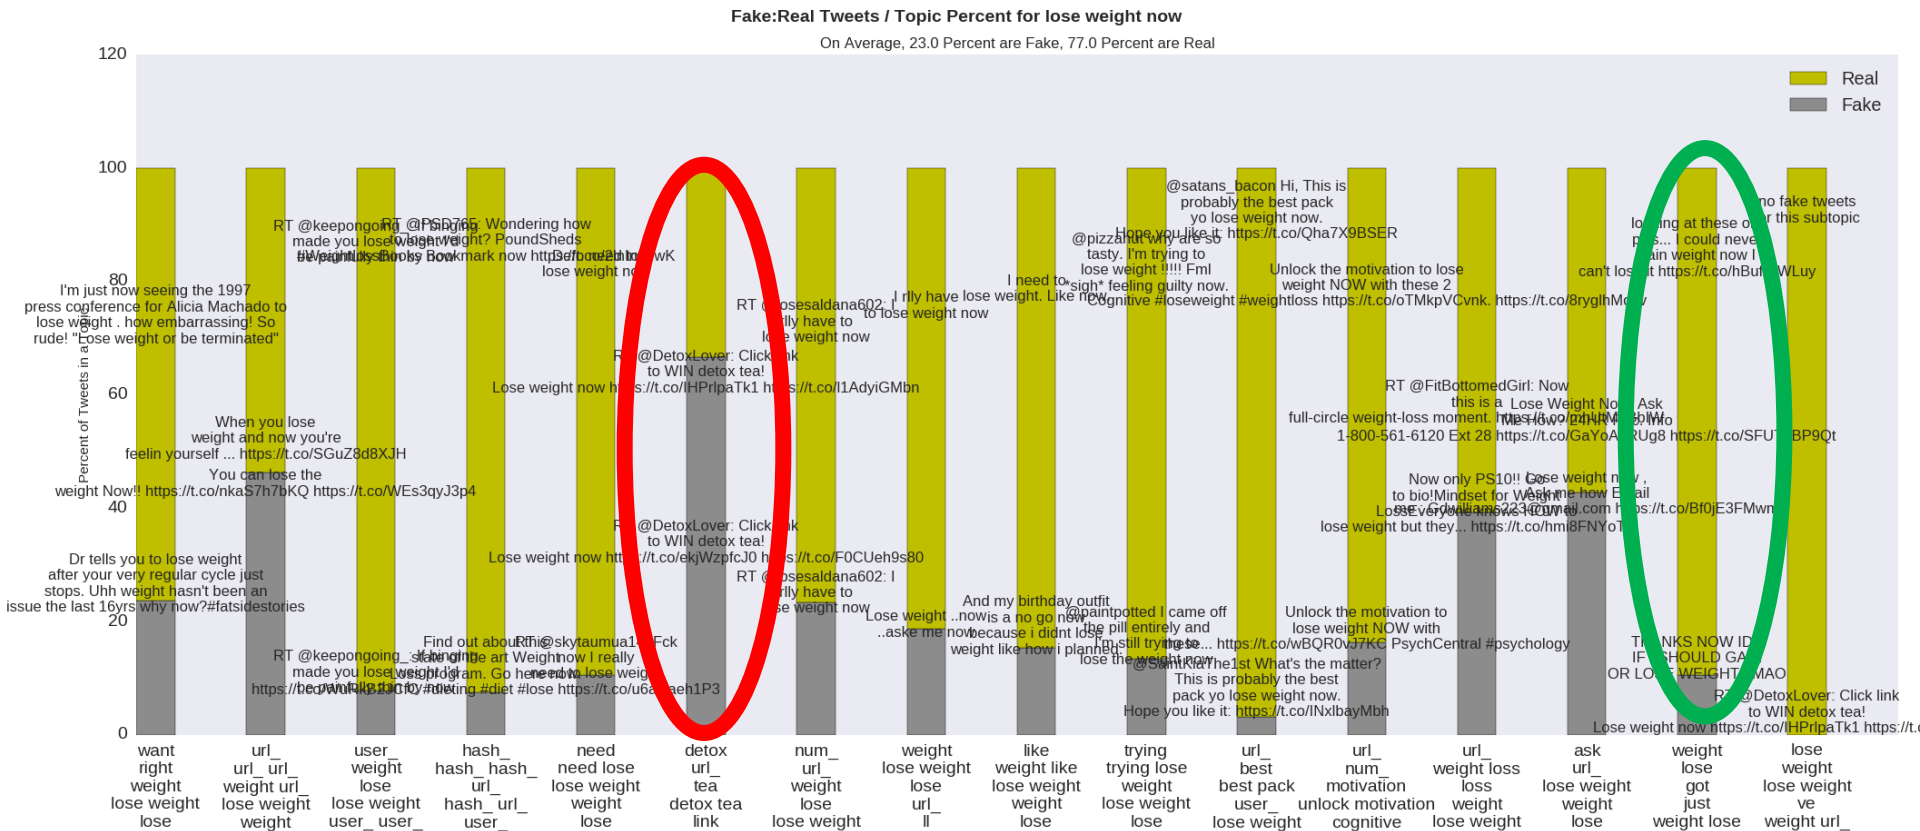
need
need lose
lose weight
weight
lose

detox
url_
tea
detox tea
link

Fake:Real Tweets / Topic Percent for lose weight now

On Average, 23.0 Percent are Fake, 77.0 Percent are Real

BotBoosted



RT @DetoxLover: Click link to WIN detox tea!
w https://t.co/ekjWzpfCJ0 https://t.co/

looking at these old pics... I could never gain weight now I can't lose it https://t.co/hBufwjWLuy

We can see which conversations are boosted by bots

Limitations & Future Work

- Twitter is very aggressive at suspending spammers
(77% suspended 1 day after 1st tweet)
(BotBoosted detects Twitter's False Negatives)
- Twitter's FREE API is not a statistically representative sample → Garden Hose API
- Strengthen the model by continuously training it with twitter's false negatives
- Benchmark Incremental "Pareto" NMF vs HDP and LDA to further improve the heuristic algorithm
- Improve corpus summarization with graph theory via tweet "centrality" based on word co-occurrence

BotBoosted

Explore and Understand
the Amount of Fabrication
of a Topic in Twitter

Thank You!

Major References:

- Aritter. Twitter NLP. (2016). Github repository https://github.com/aritter/twitter_nlp
- Cresci, S., Di Pietro, R., Petrocchi, M., Spognardi, A., Tesconi, M. (2015). Fame for Sale: efficient detection of fake Twitter followers. Retrieved from Cornell University Library. (arXiv:1509.04098)
- Joseph, K., Landwehr, P., Carley, K. (2014). Two 1% don't make a whole: Comparing simultaneous samples from Twitter's Streaming API.
- Karambelkar, B. (2015, Jan 5). How to use Twitter's Search REST API most effectively. Retrieved from <http://www.karambelkar.info/2015/01/how-to-use-twitters-search-rest-api-most-effectively/>
- Kharde, V., Sonawane, S. (2016). Sentiment Analysis of Twitter Data: A Survey of Techniques. International Journal of Computer Applications. Volum 139, No 11, April 2016.
- Kontaxis, G., Polakis, I., Ioannidis, S., Markatos, E. (2011). Detecting Social Network Profile Cloning. Retrieved from SysSec Consortium. (n.d.).
- Mori, T., Kikuchi, M., Yoshia K. (2001). Term Weighting Method based on Information Gain Ratio for Summarizing Documents retrieved by IR systems. Retrieved from <http://research.nii.ac.jp/ntcir/workshop/OnlineProceedings2/morisumm.pdf>
- Thomas, K., Grier, C., Paxson, V., Song, D. (2011). Suspended Accounts in Retrospect: An Analysis of Twitter Spam. Retrieved from <http://www.icir.org/vern/papers/twitter-susp-accounts.imc2011.pdf>

Brian Balagot

brian.balagot@gmail.com

<https://www.linkedin.com/in/briancbalagot>

<https://github.com/brityboy>

